

P2-Adversarial Search

[New Attempt](#)

Due Mar 8 by 11pm **Points** 100 **Submitting** a file upload **File Types** zip

The code base for this project: [P2_game.zip](#) 

(https://canvas.ou.edu/courses/227140/files/51261599/download?download_frd=1)

Note: you can also find this project description in the zip file provided for this project.

Introduction*

In this project, you will design Pacman and ghosts in maze environments. you will implement both minimax and expectimax search with designing evaluation function.

Submission Instruction:

Similar to project 1, create a **P2_supplement.txt** that includes the name of yours and partners (in any) in your group. Also your supplement file should explain the evaluation function you choose for Question 6 and the reason you have chosen (i.e. explain why you think this evaluation function is reasonable). **In addition please see the attached rubric for any other information that needs to be detailed for your submission.**

The zip file should follow the format of P2_student1 lastname_student2 lastname.zip

Important note: your code should be very well documented to get the full credit. The code should pass all the autograder. Also it shows the pacman performance graphically. However, we also review the code to evaluate the correctness of the code and the final grade will be based on that

Academic Misconduct: we are strict in any dishonesty and we check your code against other submissions in the class for any redundancy. If we find that the code is copied and it is not your own work, we will follow the university policy for misconduct. Please make sure you present your own work. You can still discuss with other students about the project, but you have to submit your own work and do not copy the solution from any other resources except your teammate's.

Need help? please do not hesitate to reach out the course staff if you have any question about the project during office hours. If you cannot meet during office hours, let us know and we try our best to arrange time. We are here to help you, and we want to make sure this project brings you fun along with success in understanding the course material.

The last, but not least! please, please start early to have more resource for help and you can learn more!

Good luck and have fun!

Autograder:

As in project 1, this project includes an autograder for you to grade your answers on your machine. This can be run on all questions with the command:

```
python autograder.py
```

Note: If your `python` refers to Python 2.7, you may need to invoke `python3 autograder.py` (and similarly for all subsequent Python invocations) or create a conda environment as described in project 0.

It can be run for one particular question, such as q2, by:

```
python autograder.py -q q2
```

It can be run for one particular test by commands of the form:

```
python autograder.py -t test_cases/q2/0-small-tree
```

By default, the autograder displays graphics with the `-t` option, but doesn't with the `-q` option. You can force graphics by using the `--graphics` flag, or force no graphics by using the `--no-graphics` flag.

See the autograder tutorial in Project 0 for more information about using the autograder.

Codebase:

The code for this project contains the following files:

Files you'll edit:	
<code>multiAgents.py</code>	Where all of your multi-agent search agents will reside.
Files you might want to look at:	
<code>pacman.py</code>	The main file that runs Pacman games. This file also describes a Pacman GameState type, which you will use extensively in this project.
<code>game.py</code>	The logic behind how the Pacman world works. This file describes several supporting types like AgentState, Agent, Direction, and Grid.
<code>util.py</code>	Useful data structures for implementing search algorithms. You don't need to use these for this project, but may find other functions defined here to be useful.
Supporting files you can ignore:	
<code>graphicsDisplay.py</code>	Graphics for Pacman

<code>graphicsUtils.py</code>	Support for Pacman graphics
<code>textDisplay.py</code>	ASCII graphics for Pacman
<code>ghostAgents.py</code>	Agents to control ghosts
<code>keyboardAgents.py</code>	Keyboard interfaces to control Pacman
<code>layout.py</code>	Code for reading layout files and storing their contents
<code>autograder.py</code>	Project autograder
<code>testParser.py</code>	Parses autograder test and solution files
<code>testClasses.py</code>	General autograding test classes
<code>test_cases/</code>	Directory containing the test cases for each question
<code>multiagentTestClasses.py</code>	Project 2 specific autograding test classes

Files to Edit and Submit: You will fill in portions of `multiAgents.py` during the assignment.

Note: The code base has not changed much from the previous project, but please start with a fresh installation, rather than intermingling files from project 1.

Welcome to Multi-Agent Pacman!

First, play a game of classic Pacman by running the following command:

```
python pacman.py
```

and using the arrow keys to move. Now, run the provided `ReflexAgent` in `multiAgents.py`

```
python pacman.py -p ReflexAgent
```

Note that it plays quite poorly even on simple layouts:

```
python pacman.py -p ReflexAgent -l testClassic
```

Inspect its code (in `multiAgents.py`) and make sure you understand what it's doing.

Question 1 (16 points): Reflex Agent

Improve the `ReflexAgent` in `multiAgents.py` to play respectably. The provided reflex agent code provides some helpful examples of methods that query the `GameState` for information. A capable reflex agent will have to consider both food locations and ghost locations to perform well. Your agent should easily and reliably clear the `testClassic` layout:

```
python pacman.py -p ReflexAgent -l testClassic
```

Try out your reflex agent on the default `mediumClassic` layout with one ghost or two (and animation off to speed up the display):

```
python pacman.py --frameTime 0 -p ReflexAgent -k 1
```

```
python pacman.py --frameTime 0 -p ReflexAgent -k 2
```

How does your agent fare? It will likely often die with 2 ghosts on the default board, unless your evaluation function is quite good.

Note: Remember that `newFood` has the function `asList()`

Note: As features, try the reciprocal of important values (such as distance to food) rather than just the values themselves.

Note: The evaluation function you're writing is evaluating state-action pairs; in later parts of the project, you'll be evaluating states.

Note: You may find it useful to view the internal contents of various objects for debugging. You can do this by printing the objects' string representations. For example, you can print `newGhostStates` with `print(newGhostStates)`.

Options: Default ghosts are random; you can also play for fun with slightly smarter directional ghosts using `-g DirectionalGhost`. If the randomness is preventing you from telling whether your agent is improving, you can use `-f` to run with a fixed random seed (same random choices every game). You can also play multiple games in a row with `-n`. Turn off graphics with `-q` to run lots of games quickly.

Grading: We will run your agent on the `openClassic` layout 10 times. You will receive 0 points if your agent times out, or never wins. You will receive 1 point if your agent wins at least 5 times, or 2 points if your agent wins all 10 games. You will receive an addition 1 point if your agent's average score is greater than 500, or 2 points if it is greater than 1000. You can try your agent out under these conditions with

```
python autograder.py -q q1
```

To run it without graphics, use:

```
python autograder.py -q q1 --no-graphics
```

Don't spend too much time on this question, though, as the meat of the project lies ahead.

Q1.1 Please detail in 3-4 sentences your thought process on how you planned to tackle the ReflexAgent problem in your supplementary file.

Question 2 (20 points): Minimax

Now you will write an adversarial search agent in the provided `MinimaxAgent` class stub in `multiAgents.py`. Your minimax agent should work with any number of ghosts, so you'll have to write an algorithm that is slightly more general than what you've previously seen in lecture. In particular, your minimax tree will have multiple min layers (one for each ghost) for every max layer.

Your code should also expand the game tree to an arbitrary depth. Score the leaves of your minimax tree with the supplied `self.evaluationFunction`, which defaults to `scoreEvaluationFunction`. `MinimaxAgent` extends `MultiAgentSearchAgent`, which gives access to `self.depth` and `self.evaluationFunction`. Make sure your minimax code makes reference to these two variables where appropriate as these variables are populated in response to command line options.

Important: A single search ply is considered to be one Pacman move and all the ghosts' responses, so depth 2 search will involve Pacman and each ghost moving two times.

Grading: We will be checking your code to determine whether it explores the correct number of game states. This is the only reliable way to detect some very subtle bugs in implementations of minimax. As a result, the autograder will be very picky about how many times you call `GameState.generateSuccessor`. If you call it any more or less than necessary, the autograder will complain. To test and debug your code, run

```
python autograder.py -q q2
```

This will show what your algorithm does on a number of small trees, as well as a pacman game. To run it without graphics, use:

```
python autograder.py -q q2 --no-graphics
```

Hints and Observations

- Hint: Implement the algorithm recursively using helper function(s).
- The correct implementation of minimax will lead to Pacman losing the game in some tests. This is not a problem: as it is correct behavior, it will pass the tests.
- The evaluation function for the Pacman test in this part is already written (`self.evaluationFunction`). You shouldn't change this function, but recognize that now we're evaluating *states* rather than actions, as we were for the reflex agent. Look-ahead agents evaluate future states whereas reflex agents evaluate actions from the current state.
- The minimax values of the initial state in the `minimaxClassic` layout are 9, 8, 7, -492 for depths 1, 2, 3 and 4 respectively. Note that your minimax agent will often win (665/1000 games for us) despite the dire prediction of depth 4 minimax.

```
python pacman.py -p MinimaxAgent -l minimaxClassic -a depth=4
```

- Pacman is always agent 0, and the agents move in order of increasing agent index.

- All states in minimax should be `GameStates`, either passed in to `getAction` or generated via `GameState.generateSuccessor`. In this project, you will not be abstracting to simplified states.
- On larger boards such as `openClassic` and `mediumClassic` (the default), you'll find Pacman to be good at not dying, but quite bad at winning. He'll often thrash around without making progress. He might even thrash around right next to a dot without eating it because he doesn't know where he'd go after eating that dot. Don't worry if you see this behavior, question 5 will clean up all of these issues.
- When Pacman believes that his death is unavoidable, he will try to end the game as soon as possible because of the constant penalty for living. Sometimes, this is the wrong thing to do with random ghosts, but minimax agents always assume the worst:

```
python pacman.py -p MinimaxAgent -l trappedClassic -a depth=3
```

Q2.1 Please give a short 2-3 sentence reason why Pacman always rushes the closest ghost in this case.

Question 3 (20 points): Alpha-Beta Pruning

Make a new agent that uses alpha-beta pruning to more efficiently explore the minimax tree, in `AlphaBetaAgent`. Again, your algorithm will be slightly more general than the pseudocode from lecture, so part of the challenge is to extend the alpha-beta pruning logic appropriately to multiple minimizer agents.

You should see a speed-up (perhaps depth 3 alpha-beta will run as fast as depth 2 minimax). Ideally, depth 3 on `smallClassic` should run in just a few seconds per move or faster.

```
python pacman.py -p AlphaBetaAgent -a depth=3 -l smallClassic
```

The `AlphaBetaAgent` minimax values should be identical to the `MinimaxAgent` minimax values, although the actions it selects can vary because of different tie-breaking behavior. Again, the minimax values of the initial state in the `minimaxClassic` layout are 9, 8, 7 and -492 for depths 1, 2, 3 and 4 respectively.

Grading: Because we check your code to determine whether it explores the correct number of states, it is important that you perform alpha-beta pruning without reordering children. In other words, successor states should always be processed in the order returned by `GameState.getLegalActions`. Again, do not call `GameState.generateSuccessor` more than necessary.

You must not prune on equality in order to match the set of states explored by our autograder. (Indeed, alternatively, but incompatible with our autograder, would be to also allow for pruning on equality and invoke alpha-beta once on each child of the root node, but this will not match the autograder.)

You should implement alpha-beta algorithm discussed in the class for this question (please see the pseudo-code in the related lecture).

To test and debug your code, run

```
python autograder.py -q q3
```

This will show what your algorithm does on a number of small trees, as well as a pacman game. To run it without graphics, use:

```
python autograder.py -q q3 --no-graphics
```

The correct implementation of alpha-beta pruning will lead to Pacman losing some of the tests. This is not a problem: as it is correct behavior, it will pass the tests.

Q3.1 Run the following command

```
python pacman.py -p AlphaBetaAgent -a depth=3 -l smallClassic
```

using depths 1, 2, 3, and 4. For each depth, give a short sentence describing how the agent played, did it win or lose, and if there were any noticeable patterns.

Question 4 (20 points): Expectimax

Minimax and alpha-beta are great, but they both assume that you are playing against an adversary who makes optimal decisions. As anyone who has ever won tic-tac-toe can tell you, this is not always the case. In this question you will implement the `ExpectimaxAgent`, which is useful for modeling probabilistic behavior of agents who may make suboptimal choices.

As with the search and constraint satisfaction problems covered so far in this class, the beauty of these algorithms is their general applicability. To expedite your own development, we've supplied some test cases based on generic trees. You can debug your implementation on small the game trees using the command:

```
python autograder.py -q q4
```

Debugging on these small and manageable test cases is recommended and will help you to find bugs quickly.

Once your algorithm is working on small trees, you can observe its success in Pacman. Random ghosts are of course not optimal minimax agents, and so modeling them with minimax search may not be appropriate. `ExpectimaxAgent`, will no longer take the min over all ghost actions, but the expectation according to your agent's model of how the ghosts act. To simplify your code, assume you will only be running against an adversary which chooses amongst their `getLegalActions` uniformly at random.

To see how the `ExpectimaxAgent` behaves in Pacman, run:

```
python pacman.py -p ExpectimaxAgent -l minimaxClassic -a depth=3
```

You should now observe a more cavalier approach in close quarters with ghosts. In particular, if Pacman perceives that he could be trapped but might escape to grab a few more pieces of food, he'll at least try. Investigate the results of these two scenarios:

```
python pacman.py -p AlphaBetaAgent -l trappedClassic -a depth=3 -q -n 10
```

```
python pacman.py -p ExpectimaxAgent -l trappedClassic -a depth=3 -q -n 10
```

You should find that your `ExpectimaxAgent` wins about half the time, while your `AlphaBetaAgent` always loses. Make sure you understand why the behavior here differs from the minimax case.

The correct implementation of expectimax will lead to Pacman losing some of the tests. This is not a problem: as it is correct behavior, it will pass the tests.

Q4.1 Run the AlphaBetaAgent on the following environment: `python pacman.py -p AlphaBetaAgent -l trappedClassic -a depth=3 -q -n 10`

Now run the ExpectimaxAgent on the same environment: `python pacman.py -p ExpectimaxAgent -l trappedClassic -a depth=3 -q -n 10`

Now give a 2-3 sentences detailing why the results are expected between the two algorithms.

Question 5 (24 points): Evaluation Function

Write a better evaluation function for pacman in the provided function `betterEvaluationFunction`. The evaluation function should evaluate states, rather than actions like your reflex agent evaluation function did. With depth 2 search, your evaluation function should clear the `smallClassic` layout with one random ghost more than half the time and still run at a reasonable rate (to get full credit, Pacman should be averaging around 1000 points when he's winning).

Grading: the autograder will run your agent on the smallClassic layout 10 times. We will assign points to your evaluation function in the following way:

- If you win at least once without timing out the autograder, you receive 1 points. Any agent not satisfying these criteria will receive 0 points.
- +1 for winning at least 5 times, +3 for winning all 10 times
- +1 for an average score of at least 500, +2 for an average score of at least 1000 (including scores on lost games)
- The additional points for average score and computation time will only be awarded if you win at least 5 times.

You can try your agent out under these conditions with

```
python autograder.py -q q5
```

To run it without graphics, use:

```
python autograder.py -q q5 --no-graphics
```

Q5.1 Please detail in 5-6 sentences your thought process on how you planned to tackle the betterEvaluationFunction problem.

* This project is based on the project designed in ai.berkeley.edu

Project2 Rubric

Criteria	Ratings						Pts
Question 1 AutoGrader Grading scheme: < 500: 0 points >= 500: 1 points >= 1000: 2 points ----- 10 games not timed out (0 of 0 points) Grading scheme: < 10: fail >= 10: 0 points ----- 0 wins (fail: below minimum value 1) Grading scheme: < 1: fail >= 1: 0 points >= 5: 1 points >= 10: 2 points	12 pts Full Marks	8 pts missed	4 pts missed 2 points	0 pts No Marks			12 pts
Question 1 Comments/Documentation	2 pts Full Marks			0 pts No Marks			2 pts
Question 1 Description/Justification Please detail in 3-4 sentences your thought process on how you planned to tackle the ReflexAgent problem	2 pts Full Marks			0 pts No Marks			2 pts
Question 2 AutoGrader	12 pts Full Marks	10 pts Failed 1 Test Case	8 pts Failed 2 Test Cases	5 pts Failed 3 Test Case	3 pts Failed 4 Test Cases	0 pts No Marks	12 pts
Question 2 Comments/Documentation	2 pts Full Marks			0 pts No Marks			2 pts

Criteria	Ratings						Pts
<p>Question 2 Description/Justification</p> <p>When Pacman believes that his death is unavoidable, he will try to end the game as soon as possible because of the constant penalty for living. Sometimes, this is the wrong thing to do with random ghosts, but minimax agents always assume the worst. Please run the following command, and please give a short 2-3 sentence reason why Pacman always rushes the closest ghost in this case:</p> <p>python pacman.py -p MinimaxAgent -l trappedClassic -a depth=3</p>	<p>2 pts Full Marks</p>	<p>0 pts No Marks</p>					2 pts
<p>Question 2 Pacman Testing</p> <p>python pacman.py -p MinimaxAgent -l minimaxClassic -a depth=4</p>	<p>4 pts Full Marks</p>	<p>2 pts Failed more than 12/18 attempts</p>		<p>0 pts Failed more than 15/18 attempts</p>			4 pts
Question 3 AutoGrader	<p>14 pts Full Marks</p>	<p>11.67 pts Failed 1 Test Case</p>	<p>9.33 pts Failed 2 Test Cases</p>	<p>5.83 pts Failed 3 Test Case</p>	<p>3.5 pts Failed 4 Test Cases</p>	<p>0 pts No Marks</p>	14 pts
<p>Question 3 Comments/Documentation</p>	<p>2 pts Full Marks</p>			<p>0 pts No Marks</p>			2 pts

Criteria	Ratings						Pts
Question 3 Pacman Description Run the following command python pacman.py -p AlphaBetaAgent -a depth=3 -l smallClassic using depths 1, 2, 3, and 4. For each depth, give a short sentence describing how the agent played, did it win or lose, and if there were any noticeable patterns.	4 pts Full Marks			0 pts No Marks			4 pts
Question 4 AutoGrader	12 pts Full Marks	10.01 pts Failed 1 Test Case	7.99 pts Failed 2 Test Cases	4.99 pts Failed 3 Test Case	3 pts Failed 4 Test Cases	0 pts No Marks	12 pts
Question 4 Pacman Testing : python pacman.py -p ExpectimaxAgent -l minimaxClassic -a depth=3	4 pts Full Marks	2 pts Failed more than 6/18 attempts		0 pts Failed more than 9/18 attempts			4 pts
Question 4 Comments/Documentation	2 pts Full Marks			0 pts No Marks			2 pts

Criteria	Ratings							Pts
<p>Question 4 Expectimax vs AlphaBeta</p> <p>Run the AlphaBetaAgent on the following environment: python pacman.py -p AlphaBetaAgent -l trappedClassic -a depth=3 -q -n 10</p> <p>-----</p> <p>Now run the ExpectimaxAgent on the same environment: python pacman.py -p ExpectimaxAgent -l trappedClassic -a depth=3 -q -n 10</p> <p>-----</p> <p>Now give a 2-3 sentences detailing why the results are expected between the two algorithms.</p>	<p>2 pts</p> <p>Full Marks</p>			<p>0 pts</p> <p>No Marks</p>				<p>2 pts</p>
Question 5 AutoGrader	<p>16 pts</p> <p>Full Marks</p>	<p>13.35 pts</p> <p>Failed 1 Test Case</p>	<p>10.65 pts</p> <p>Failed 2 Test Cases</p>	<p>6.65 pts</p> <p>Failed 3 Test Case</p>	<p>4 pts</p> <p>Failed 4 Test Cases</p>	<p>0 pts</p> <p>No Marks</p>	<p>16 pts</p>	
Question 5 Comments/Documentation	<p>2 pts</p> <p>Full Marks</p>			<p>0 pts</p> <p>No Marks</p>				<p>2 pts</p>
Question 5 Description/Justification Please detail in 5-6 sentences your thought process on how you planned to tackle the betterEvaluationFunction problem.	<p>6 pts</p> <p>Full Marks</p>			<p>0 pts</p> <p>No Marks</p>				<p>6 pts</p>
Total Points: 100								